

Hugh Baldwin
up2157117

Operating Systems and Internetworking

M30233

TB1

University of Portsmouth
BSc Computer Science
2nd Year

Contents

I	Operating Systems	2
1	Operating Systems Cheatsheet	3
2	Lecture - Introduction	4
3	Lecture - Concurrency	7
II	Internetworking	8
4	Lecture - Network Services	9
5	Lecture - IP Addresses and Subnetting I	12

Part I

Operating Systems

Operating Systems Cheatsheet

Term	Definition
Kernel Mode	The CPU mode in which all assembly instructions can be used. This mode is usually used by the Operating System
User Mode	The CPU mode in which some assembly instructions cannot be used, such as IN and OUT. This mode is usually used by Application Software
System Software	The operating system and related utilities that control the computer, and provide essential functionality, such as keyboard and mouse input handling
Application Software	Programs that allow the user to perform specific functions with a computer, such as word processing or internet browsing
Interrupts	A signal sent to the CPU, usually from an I/O device, that requests the attention of the CPU to handle an urgent operation

Lecture - Introduction

13:00

26/09/23

Tamer Elboghdadly

Operating Systems

- The Operating System sits inbetween the hardware and application software
- It usually is not the actual GUI, but provides functionality for the applications which implement it
- OSes typically provide abstractions for applications so that they can run on different hardware

User and Kernel Mode

- User mode
 - The programs that a user directly interacts with
 - Uses an API to access hardware, rather than having direct access
- Kernel mode
 - The programs that run the operating system
 - Has direct access to hardware

System and Application Software

- Application software
 - Programs that allow a user to perform a task
 - Requires the support of system software to run
- System software
 - Software directly related to the operating system
 - Manages the boot process
 - Hardware drivers
 - File system management

The main functions of an Operating System

- Resource management
 - Manages the memory, CPU and other hardware to allow multiple programs to run concurrently
 - Handles requests from applications to allocate more resources
- "Extended machine"
 - Handles reading and writing to control registers, handling interrupts, etc
 - Provides higher level APIs for other software to interact with the hardware

CPU Organisation

Registers

Name	Use	Description
EAX	Accumulator	The default register for many addition and multiplication instructions
EBX	Base	Stores the base address during memory addressing.
ECX	Count	The default counter for repeat (REP) prefix instructions and LOOP instructions.
EDX	Data	Used for multiply and divide operations
ESI	Source Index	Store source index
EDI	Destination Index	Store destination index
EBP	Base Pointer	Mainly helps in referencing the parameter variables passed to a subroutine.
ESP	Stack Pointer	Provides the offset value within the program stack.

Figure 2.1: General purpose registers in an x86 Intel CPU

- There are also special purpose registers, such as the Program Counter (PC) and the Program Status Word
- The Program Status Word sets the mode in which the CPU is operating

Assembly Language

- Assembly Language is the lowest-level programming language before Machine Code. It is slightly abstracted from machine code and uses neumonic symbols to represent instructions
- For example:
 - `MOV EBX, EAX`
 - This copies the value in EAX into the EBX register
 - `ADD EBX, 4`
 - This adds the value 4 to the EBX register
 - This is equivalent to $b = a + 4$
- The MOV instruction can also be used to move values between registers and the main memory
- I/O devices such as hard disks have a set of ports, which can be accessed to control the device and transfer data
- The special instructions IN and OUT are used to read from or write to ports
- For example:
 - `IN EAX, 368`
 - This copies the value from the port 368 into the EAX register
- The IN and OUT instructions can only be used when the CPU is running in Kernel Mode

User and Kernel Mode

- CPUs support running in two different modes: Kernel mode and User mode
- When running in User mode, some instructions (such as IN and OUT) cannot be used
- Typically, the Kernel of an operating system will run mostly or entirely in Kernel Mode

Interrupts

- When an external device needs to gain the attention of the CPU, it sends an interrupt
- A couple of examples are as follows
 - When a disk has requested data in it's buffer
 - When the key on an old PS/2 keyboard is pressed
- When the CPU receives an interrupt, it must abandon whatever it is currently doing and run an Interrupt Handler routine
- Interrupt Handlers run in Kernel mode
- The CPU places whatever resources it was using onto the stack and executes the interrupt handler. Once the handler has finished, it returns to the execution path stored on the stack

Lecture - Concurrency

13:00

03/10/23

Tamer Elboghdadly

- A sequential system is one where multiple threads are executed to completion, one after the other
- A concurrent system is one where multiple threads are executed at the same time
- A parallel system is one where multiple threads are executed in parallel, usually on multiple hardware threads
- There are several issues that may arise from concurrency, especially in cases where threads of the same process are run on multiple processors in parallel
 - A thread is a sequence of instructions defined by a program or part thereof
 - A process has one or more threads, each of which may be run in parallel
 - Threads within a single process can share resources with each other, but have their own thread of execution

Part II

Internetworking

Lecture - Network Services

09:00

25/09/23

Athanasios Paraskelidis

DHCP

- DHCP stands for Dynamic Host Configuration Protocol
- The main motivation of DHCP was to provide a set of configuration parameters to automatically configure new devices as they are added to the network. The main parameters are as follows:
 - IP address
 - Gateway address
 - Subnet mask
 - DNS server address
- Before DHCP, either devices were configured manually, or the bootstrap protocol (BOOTP) was used
- The main improvements over BOOTP are:
 - Support for temporary allocation of IP addresses
 - DHCP clients can automatically discover the local DHCP server
 - Once the server is setup, there is almost 0 human interaction unless something goes wrong
 - Still compatible with BOOTP clients
- A lease is the length of time that an allocated IP address can be used before either a new address is needed, or a request to continue using the current address needs to be approved
- An IP address can be released by the client if it is no longer needed, e.g. the device shuts down or no longer needs to communicate on the network
- Advantages
 - Saves manually configuring every single device
 - Ability to move to a different network without having to reset any network settings on the device
 - Allows more efficient use of the IP space, as inactive devices do not need to keep a lease on an address
- Disadvantages
 - DHCP uses UDP to configure devices, so the communication is unreliable and insecure
 - Possibly allows unauthorised clients, but this can be avoided using MAC address white/blacklisting
 - Potential for malicious DHCP servers setup on a network that provide incorrect network settings

DNS

- DNS stands for Domain Name System
- This is the system which devices use to convert the human-readable domain names into IP addresses which computers can use to communicate to the server
- DNS is a globally distributed mapping database between domain names and IP addresses
- There are 3 main components
 - A **name space**
 - **Servers** that make the name space available
 - **Resolvers** which make the request from clients to servers
- As DNS is globally distributed, some data is maintained locally, but also retrievable globally
 - No single server holds all DNS records
- A DNS lookup can be performed by any device
 - Remote DNS data is usually cached locally to improve performance
- DNS has 'loose coherency'
 - The database is always internally consistent
 - Each version of the database has a serial number, which is incremented every time the database changes
 - Changes to the master copy of the database are replicated to secondary servers regularly, depending upon the timing set by the zone administrator
 - Cached data expires depending upon timing set by the zone administrator
- There is no limit to the size of the database
 - Having a very large number of records on one server would decrease performance
 - Therefore, the database is spread across many servers around the internet
- There is no limit to the number of queries which can be made at any time, or by a single user
- Queries are usually distributed between multiple DNS servers as well as local caches
 - e.g. nameserver1 and 2
- Clients can query and use the data from any server, primary or secondary
- Clients will typically have their own local cache of more frequently accessed records
- DNS uses both TCP and UDP
 - TCP is used for communication between servers, for example when replicating records from a primary to secondary server
 - UDP is used for communication between clients and servers
- The database can be updated dynamically
 - Add, delete or modify any record on the server
 - These only need to be performed on the main server, as the secondary servers will replicate the changes over time

- There are two main types of servers
 - Authoritative
 - * Primary server - Where data is added and modified
 - * Secondary server - Servers which replicate the primary server to share the load
 - Non-authoritative
 - * Caching servers - temporarily retain records from authoritative servers to improve resolving performance and reduce load on authoritative servers
- Domains can be resolved either Iteratively or Recursively
 - Iteratively
 - * The client's domain name resolver starts by querying the root nameserver
 - * The root nameserver responds with the address of the nameserver on the next level down
 - * The domain name resolver then queries this nameserver, and so on until the nameserver with the full domain is found
 - Recursively
 - * The client's domain name resolver queries the root nameserver
 - * The root nameserver itself queries the nameserver on the next level down
 - * This process repeats until the nameserver with the full address is found
 - * The IP address of the domain is then passed back up the chain and to the requester

Domain Names

- A domain name is the sequence of labels from a node to the root, separated by dots
 - e.g. port.ac.uk has 3 labels
 - * port
 - * ac
 - * uk
 - There can be up to 127 labels in a domain name
 - But there can only be 255 characters in the domain overall
- The root domain or Top Level Domain (TLD) of a domain is the final label, e.g. the TLD of port.ac.uk would be uk
- A subdomain is any domain which resides below the TLD. In the case of port.ac.uk,
 - uk is the TLD
 - ac is the Second Level Domain
 - port is the actual domain name
- Name servers store information about domains in units called zones
 - Each zone usually corresponds to a subdomain, for example the .uk TLD has many sub-zones, such as .ac.uk, .gov.uk and .co.uk

Lecture - IP Addresses and Subnetting I

09:00

02/10/23

Athanasios Paraskelidis

- Layer 3 networking is the layer that handles routing data between networks, such as a local and wide area network
- Layer 3 functionality is spread over the network - in routers as well as client devices

Internet Protocol

- Connectionless
- Best effort delivery
 - Unreliable
 - Packets may not arrive at all, or in the wrong order
 - There is no built-in error handling other than error checking
- When transmitting
 - Encapsulates data from the transport layer in datagrams
 - Adds header data to the datagrams (source and destination addresses, time to live, etc)
 - Apply the routing algorithm once on the internet
 - Will continue to route the packet until either it reaches the target network card, or exceeds it's time to live
- When receiving
 - Check the validity of all datagrams
 - Reads the header and checks if forwarding is needed
 - * If the destination address is on the LAN, send the datagram to the destination
 - * If it isn't, relay it to the next router on it's route

IPv4 & Subnetting

- 32-bit string
 - 4*8 bit integers (1 byte, or 1 octet)
 - Each byte can be 0-255
 - Intended to be human-readable
 - Allows for 2^{32} or 4, 294, 967, 296 theoretical addresses, less in practice
- The IP address is split into two fields
 - The Network Prefix
 - The Host ID

- Any devices with the same network prefix are on one network
- Two devices cannot have the same Host ID
- The network prefix can be 8, 16 or 24 bits in length, leaving 24, 16 or 8 bits for Host IDs
- There are 3 main classes of IP addresses
 - Class A: Network prefix always starts with 0, followed by 7 bits, and 24 bits for the Host ID (1.X.X.X - 126.X.X.X)
 - Class B: Network prefix always starts with 10, followed by 14 bits, and 16 bits for the Host ID (128.0.X.X.X - 191.255.X.X)
 - Class C: Network prefix always starts with 110, followed by 21 bits, and 8 bits for the Host ID (192.0.0.X - 223.255.255.X)
- There are also 2 more reserved classes
 - Class D: Network prefix starts with 1110, followed by a 28 bit multicast ID (224.X.X.X - 239.X.X.X)
 - Class E: Network prefix starts with 1111, followed by 28 bits reserved for experimentation (240.X.X.X - 255.X.X.X)
- The Class A network 127.X.X.X is reserved for loopback addresses, every computer has 127.0.0.1 as a loopback address, some computers may have more than 1
- Each class of network has 2 less usable host addresses than would be expected, e.g.
 - A class A network provides $2^{24} - 2$ host addresses (16, 777, 214)
 - A class B network provides $2^{16} - 2$ host addresses (65, 534)
 - A class C network provides $2^8 - 2$ host addresses (254)
- This is because X.X.X.0 is the network address, which cannot be assigned to a computer, and X.X.X.255 is used as the network's broadcast address
- There are 3 ranges of IP addresses that can be used by anyone on a local network, one in class A, B and C
 - Class A: 10.0.0.0 - 10.255.255.255
 - Class B: 172.16.0.0 - 172.31.255.255
 - Class C: 192.168.0.0 - 192.168.255.255
- In order to connect to devices on the internet, a NAT (Network Address Translation) service is needed on the network (usually on the router)

Subnet Masks

- A subnet mask tells the computer which parts of the IP address are reserved for the network prefix
- The default subnet mask for each class is as follows:
 - Class A: 255.0.0.0
 - Class B: 255.255.0.0
 - Class C: 255.255.255.0
- The subnet mask also determines how many devices can be on the network at a time, as it sets how many bits are left for the Host ID
- A network address is the network prefix, with all host bits set to 0, e.g.

- The network address of 10.128.47.87 is 10.0.0.0
 - The network address of 192.168.10.104 is 192.168.10.0
- But why do we need subnets?
 - If you are assigned a Class B network address (e.g. 128.147.0.0), you have up to 65534 host addresses available
 - It's not possible to effectively manage that many clients on a single physical network, and would be a waste of most of those IP addresses
 - Therefore, we should split the network into subnets which each have a much more manageable number of devices
 - e.g. we could use the 3rd byte of the address as a subnet ID and have 256 subnets, each with 254 hosts
 - Since the subnet ID is configurable, we could also use 4 bits for the subnet ID, allowing 16 subnets with 4094 hosts each
- There is an alternative way to write subnet masks, which is called the slash notation
 - This is noted by adding the number of bits reserved for the network address after a slash at the end of the IP
 - For example, you could say that you have a Class B network with the ID 148.197.0.0 and therefore subnet mask 255.255.0.0, or you could simply write it as 148.197.0.0/16
- This allows us to use custom length subnet masks when creating subnets
- For example, you may have a Class B address and want to have 32 subnets. To do this, you would need an additional 5 bits for a subnet ID, and therefore should use the subnet mask /21. This allows for 32 subnets with $2^{11} - 2$ or 2,046 hosts each